

Shayan Shariff

shayan-ahmed.shariff6@etu.univ-lorraine.fr | [shayanshariff.github.io](https://github.com/shayanshariff) | linkedin.com/in/shayan-shariff

EDUCATION

Université de Lorraine

Master of Science in Natural Language Processing

- Thesis: Understanding Context-Memory Conflicts in Retrieval-Augmented Code Generation

09/2024 – 08/2026

Nancy, France

Habib University

Bachelor of Science in Computer Science, Minor in Comparative Literature

08/2019 – 05/2023

Karachi, Pakistan

EXPERIENCE

Research Intern

LORIA / Steerway

- Built a pipeline measuring whether code models follow corrupted API documentation or their training when the two conflict, with AST-based classification of generated code
- Conducted scaling analysis across 31 models from seven families of LLMs (0.5B–32B).
- Performed mechanistic interpretability experiments (logit lens, activation patching) identifying scale-dependent weakening of memory and context heads

March 2026 – Sept. 2026

Nancy, France

Research Intern

LORIA

- Analysis of several industry standard LLM question-answering benchmarks for stereotypical bias
- Manual annotation of stereotypical bias in the MMLU LLM benchmark
- Presented completed research as a poster at the *Forum des Sciences Cognitives & TAL 2025* — **Best Poster Award**

June 2025– Aug. 2025

Nancy, France

Software Engineer

Alt Ventures

- Built an AI interview platform (custom-prompted voice bots with function calling and data collection) using React, Node.js/Express, and the Vapi and Hume voice AI APIs
- Developed a cross-platform B2B SaaS application used by 30+ enterprise clients

July 2022 – Aug 2024

Lahore, Pakistan (Remote)

AI SAFETY TRAINING AND HACKATHONS

Frontier AI Security Training (FAST)

Participant (one of 18), Singapore AI Safety Hub

- AI control hackathon: designed and attacked monitoring protocols for an untrusted model as red and blue teams
- Open-weight security labs (abliteration, backdoor training and detection, distillation), AI verification, wargaming

Sept. – Oct. 2026

Singapore

Secret Loyalties Hackathon — Apart Research (top 10%)

Jul. 2026

Technical AI Safety Course — BlueDot Impact

Jul. 2026

PUBLICATIONS

Are Benchmarks a Source of Bias in LLMs? MMLU, CommonSenseQA and MGSM under the Microscope. Fanny Duce, **Shayan Ahmed Shariff**, Ibrahim Al Kotob, Binesh Arakkal Remesh, Lucie Digoin-Caparras, Aurélie Névéol, Karën Fort. *EMNLP 2026 Main (CORE A*)*, Oral, Budapest, 2026.

Les benchmarks sont une source de biais des LLM : MMLU, CommonSenseQA et MGSM au microscope. Fanny Duce, Lucie Digoin-Caparras, Ibrahim Al Kotob, **Shayan Ahmed Shariff**, Binesh Arakkal Remesh, Aurélie Névéol, Karën Fort. *Actes de la 33e Conférence sur le Traitement Automatique des Langues Naturelles (TALN)*, 2026.

Bridging Language and Culture: Applying NLP Methods for Topic Classification of Couplets using BERT. **Shayan Shariff**, Ali Adnan, and Muhammad Qasim Pasta. ICMLC 2024, Shenzhen, China. ACM pp. 582–586.

SKILLS

ML & NLP: PyTorch, HuggingFace (Transformers, Datasets), RAG systems, Code Generation

Engineering: Python, JavaScript, React, Node.js/Express, SQL, REST APIs, Git, Grid5000

Languages: Urdu (Native), English (Fluent), French (B1)